




SEO → Search Engine Optimization.

Let's { "Wireless blue tooth headphones"
→ SEO Summary

You are an assistant that summarizes products for SEO and extracts key features. Return JSON with the keys, summary, and the features.

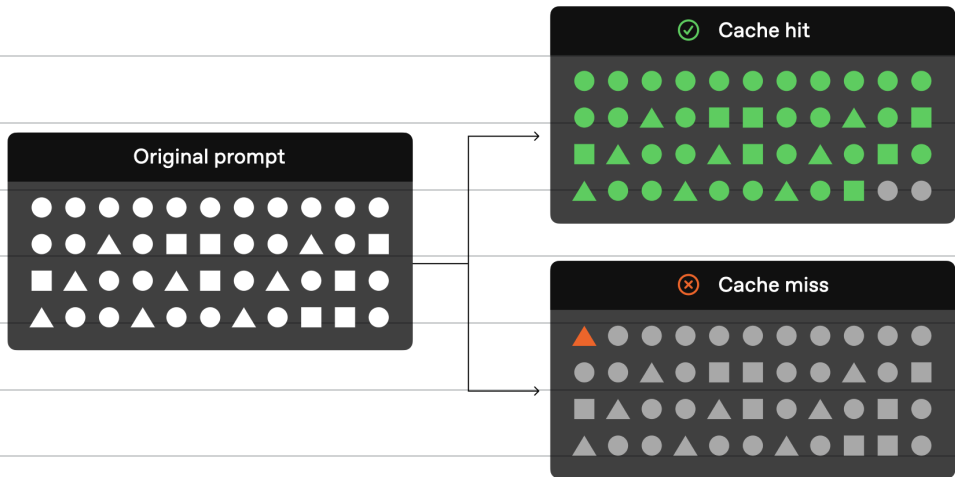
Monitoring Cache Usage

API calls to supported models will automatically benefit from Prompt Caching on prompts longer than 1,024 tokens. The API caches the longest prefix of a prompt that has been previously computed, starting at 1,024 tokens and increasing in 128-token increments. If you reuse prompts with common prefixes, we will automatically apply the Prompt Caching discount without requiring you to make any changes to your API integration.

System prompt 1100 tokens (1024 tokens cached
76 tokens not cached
1300 tokens (1024 tokens - 1024
+ 128 tokens - 1152
+ 128 tokens - 1280
20 tokens not cached

1500 tokens

System prompt + "Pls write on opti SEO
for { }"



Batching

1. When you want to process 1000s of results, and don't need the result instantly.
2. You upload relevant file/prompt when you start batch job.
3. You poll, until it's done.
4. Download the results at once.

Client Server



Queue your request.

Text tokens Per 1M tokens · Batch API price ☐

Input \$0.05	Cached input \$0.005	Output \$0.40
------------------------	--------------------------------	-------------------------

Quick comparison Input Cached input Output

Text tokens Per 1M tokens · Batch API price ☒

Input \$0.025	Cached input \$0.003	Output \$0.20
-------------------------	--------------------------------	-------------------------

Can. you cover a bit of MLOps/AIOPS/AgentOps